


The LLM judge is ungoverned and model-dependent; RAILS is not

A. Defective work, self-report only
(evidence below the floor)



B. Legitimate work
(throughput / specificity)

