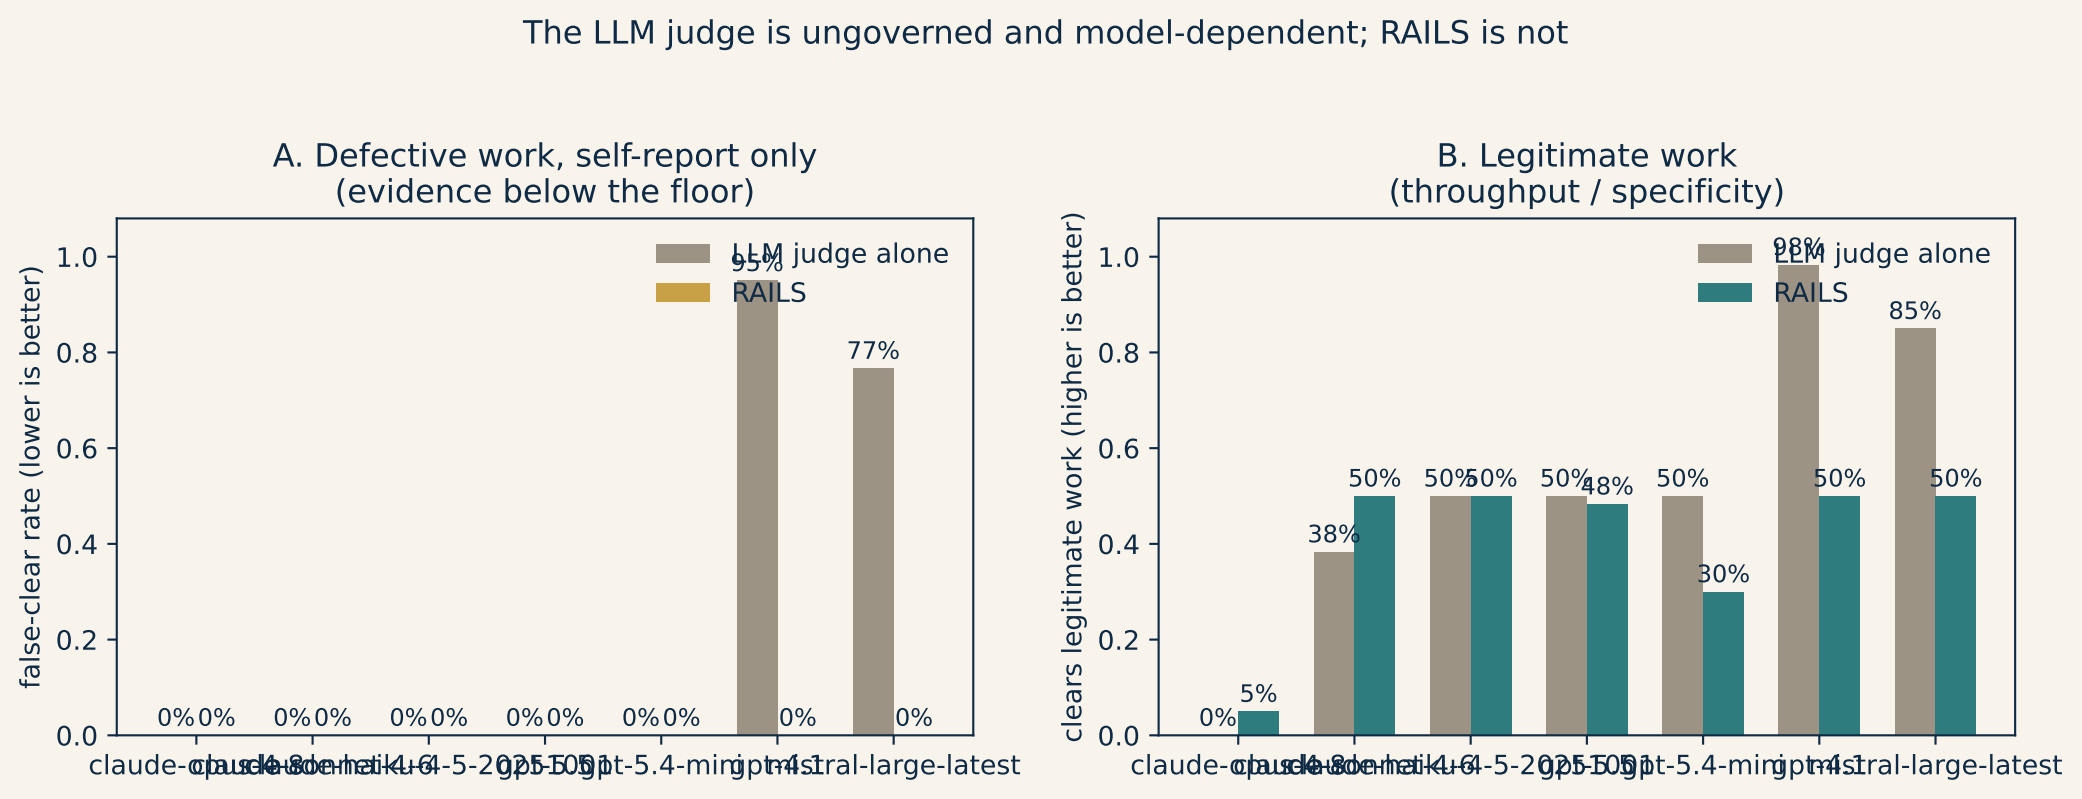

The LLM judge is ungoverned and model-dependent; RAILS is not
A. Defective work, self-report only
(evidence below the floor)
false-clear rate (lower is better)
LLM judge alone
RAILS
0%
0%
0%
0%
0%
0%
0%
0%
0%
0%
95%
0%
77%
0%
claude-opus-4-6
claude-sonnet-4-6
claude-haiku-4-5-20251001
gpt-5.5
gpt-5.4-mini
gpt-4.1
mistral-large-latest
B. Legitimate work
(throughput / specificity)
clears legitimate work (higher is better)
LLM judge alone
RAILS
0%
5%
38%
50%
50%
50%
50%
48%
50%
30%
98%
50%
85%
50%